

PFind: Parallel File Inspection and Node Discovery

Denis Koshelev, Aditya Kaushik, Ryan Lynch

Georgia Institute of Technology

April 23, 2024

Agenda

1. Problem statement

- Limitations of the current `find` command

Agenda

1. Problem statement
 - Limitations of the current `find` command
2. Our parallel implementation
 - How `pfind` works
 - Approach 1: Recursive version
 - Approach 2: Iterative BFS

Agenda

1. Problem statement
 - Limitations of the current `find` command
2. Our parallel implementation
 - How `pfind` works
 - Approach 1: Recursive version
 - Approach 2: Iterative BFS
3. Current progress

Agenda

1. Problem statement
 - Limitations of the current `find` command
2. Our parallel implementation
 - How `pfind` works
 - Approach 1: Recursive version
 - Approach 2: Iterative BFS
3. Current progress
4. Evaluation
 - Data collection and challenges
 - Benchmarking with competitors

Limitations of the current **find** command


```
koselev@lawn-128-61-61-164:~/why_find_sucks
> tree
.
├── 1_sequential_search
│   ├── bad.txt
│   └── terrible.txt
├── 2_no_multicore_support
│   ├── how_dare.txt
│   └── its_2024.zip
└── 3_inefficient_in_large_systems
    └── it_takes_forever.txt

4 directories, 5 files


> find . -name "it_takes_forever.txt"
./3_inefficient_in_large_systems/it_takes_forever.txt

>
```

How **pfind** works



How **pfind** works



How **pfind** works

Our main ideas for parallel implementation are:

How **pfind** works

Our main ideas for parallel implementation are:

- Using **directories** as parallelizable units of work

How **pfind** works

Our main ideas for parallel implementation are:

- Using **directories** as parallelizable units of work
- Using **heuristics** such as directory size and/or number of files to split up work amongst threads

Approach 1: Recursive approach

An outer function calls a recursive function `processDirectory()` that takes start folder and target to find as parameters:


```
function parallelFind(root, target) {  
    #pragma omp parallel single  
    processDirectory(root, target)  
}
```

Approach 1: Recursive approach

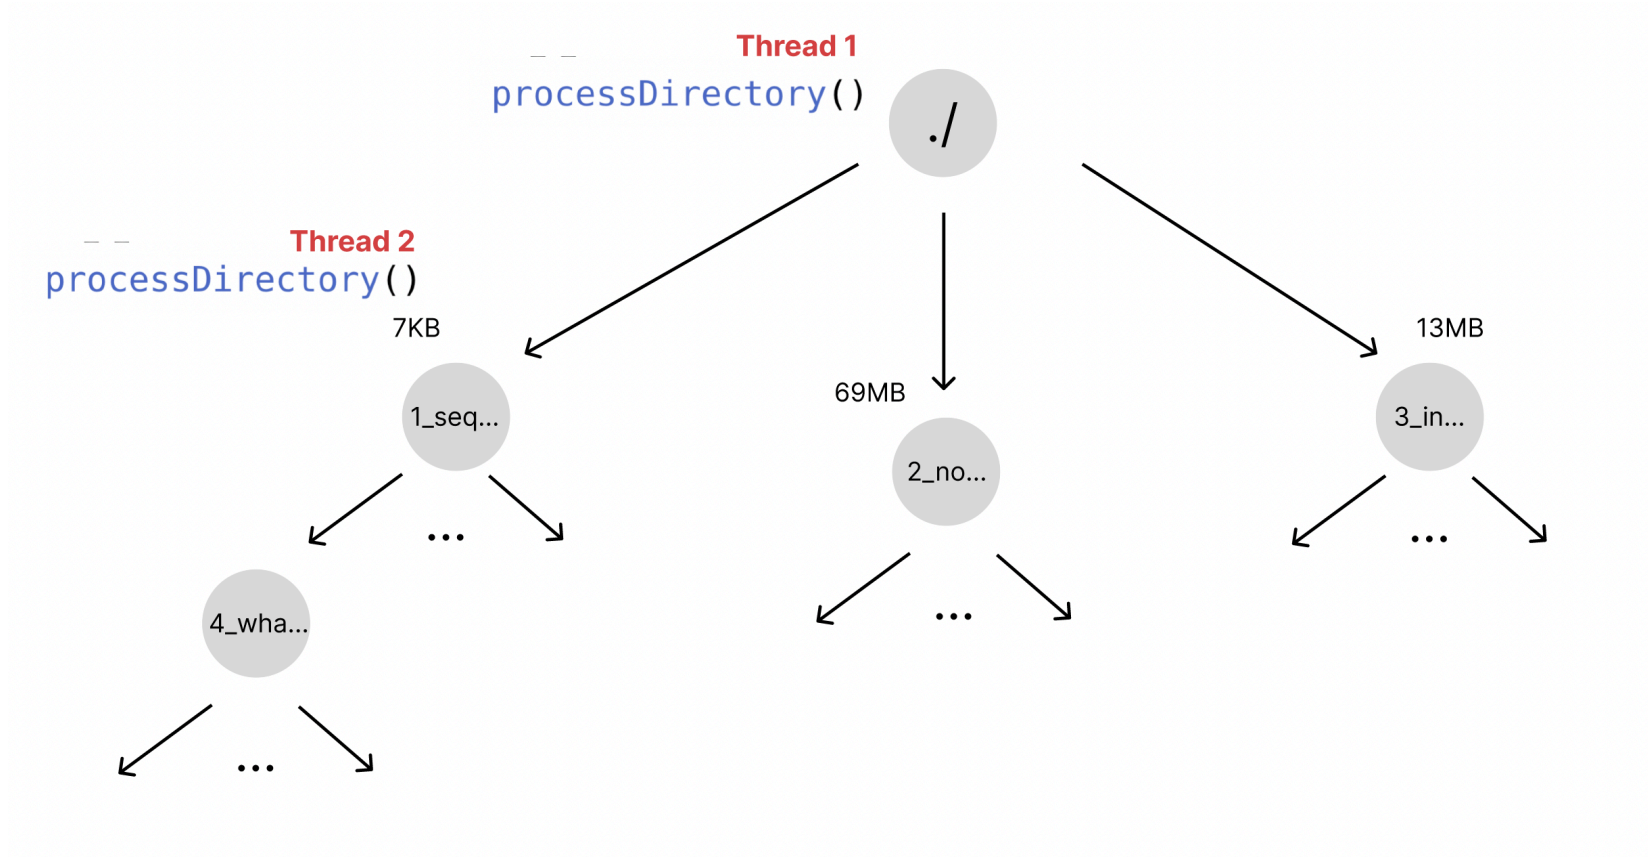
`processDirectory()` function recursively makes parallel calls to itself whenever new directory is encountered or prints path if file found:

```
function processDirectory(directory, target) {  
    for entry in directory:  
        if entry is directory:  
            #pragma omp task  
            processDirectory(entry, target)  
        else if entry matches target:  
            print(entry)  
}
```

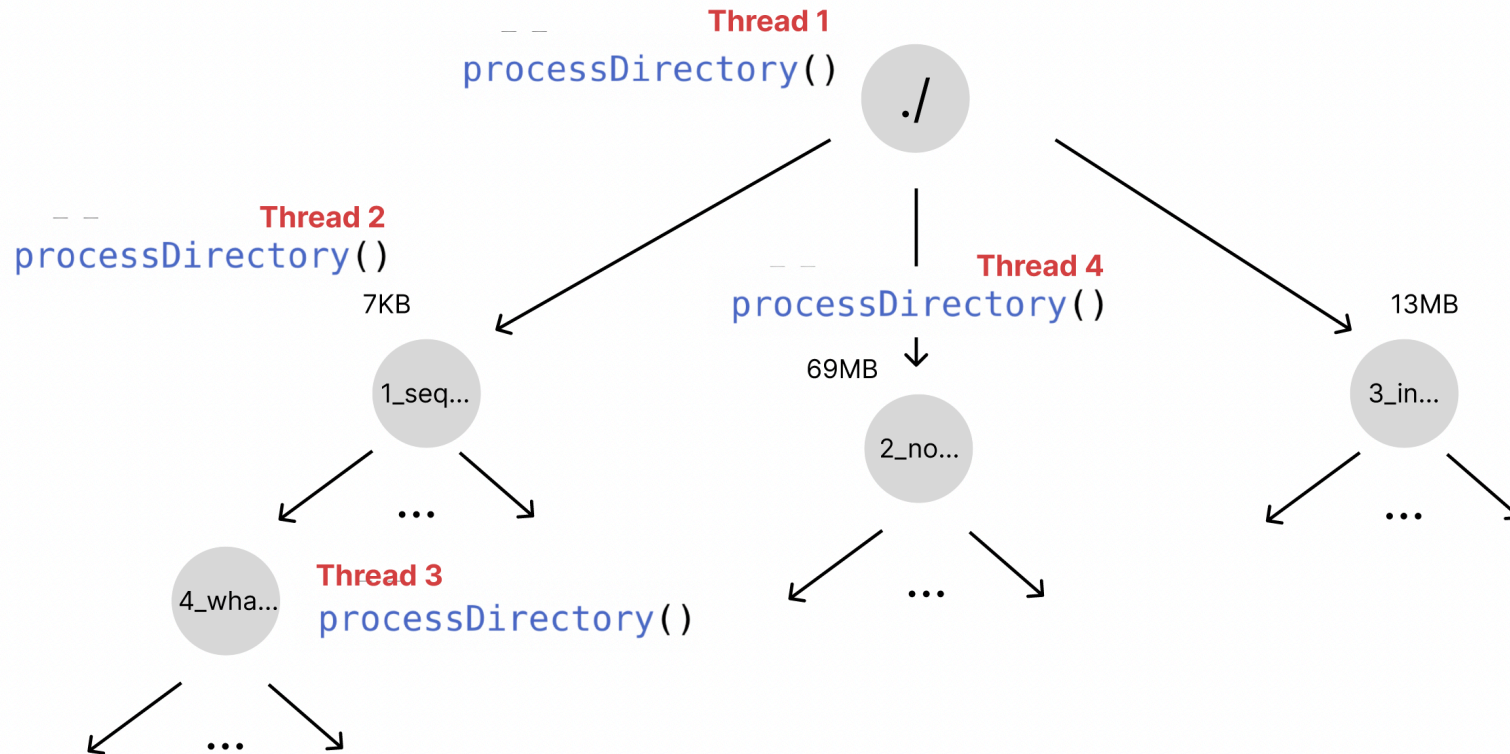
Approach 1: Recursive approach



Approach 1: Recursive approach



Approach 1: Recursive approach



Approach 2: BFS

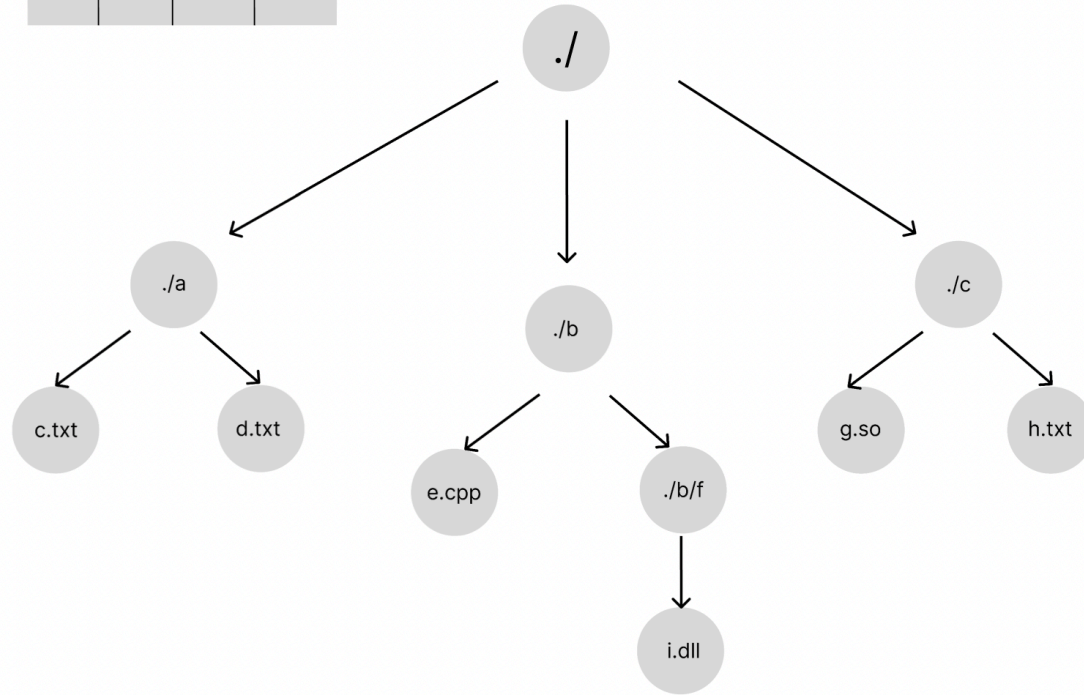
```
omp_set_lock(&queue_lock);  
if (!directory_queue.empty()) {  
    current_directory = directory_queue.front();  
    directory_queue.pop();  
    active_tasks++;  
}  
omp_unset_lock(&queue_lock);
```

Approach 2: BFS


```
omp_set_lock(&queue_lock);  
directory_queue.push(full_path);  
omp_unset_lock(&queue_lock);` )
```

Approach 2: BFS


Queue



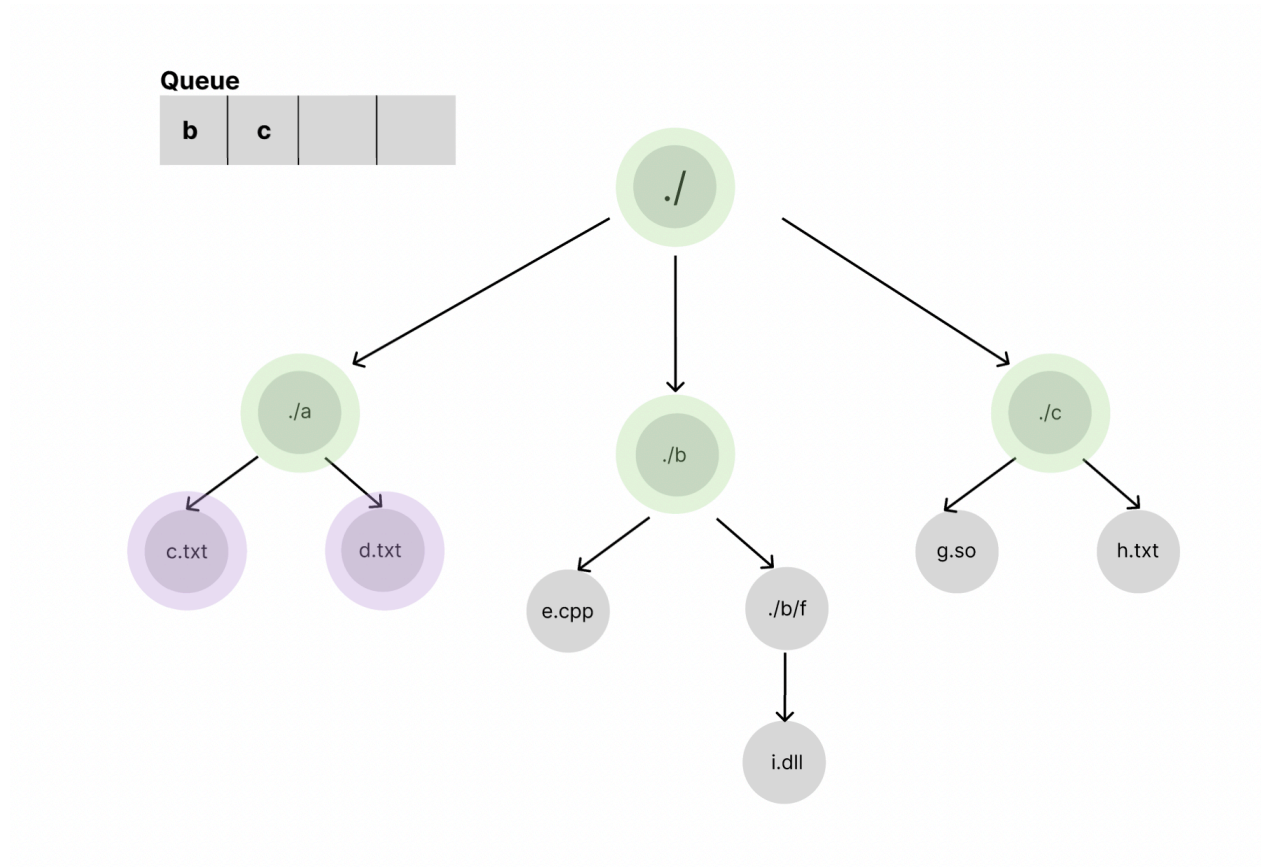
Approach 2: BFS



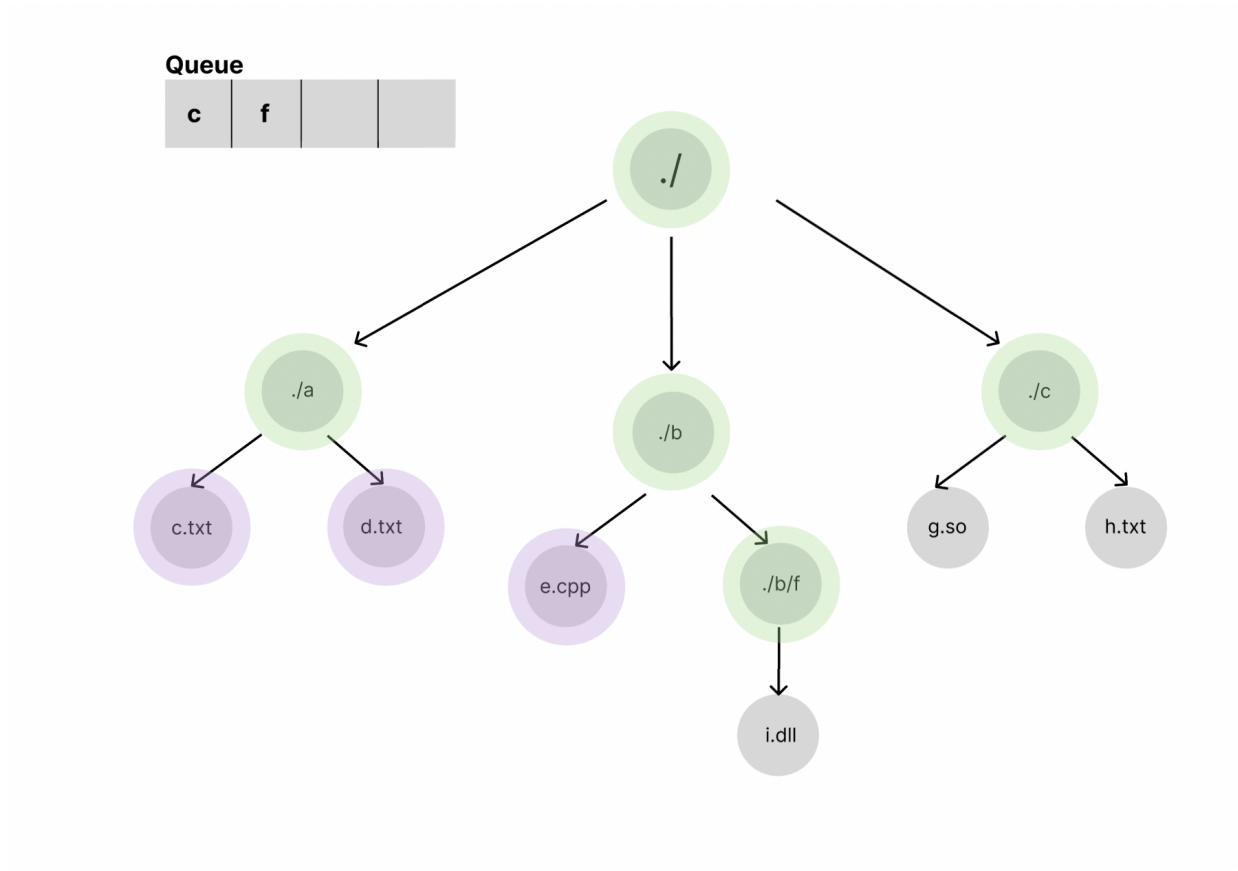
Approach 2: BFS



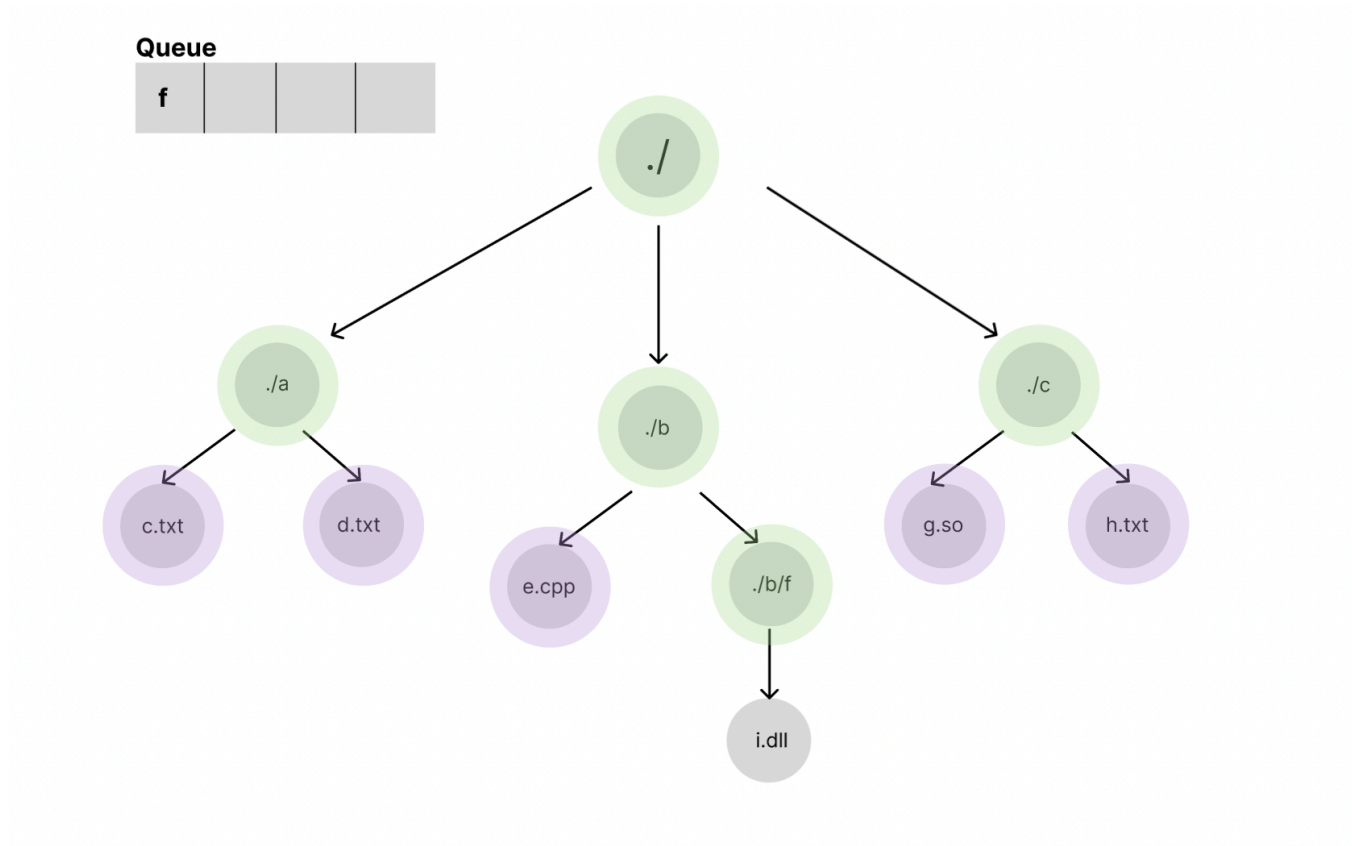
Approach 2: BFS




Approach 2: BFS



Approach 2: BFS



Approach 2: BFS



Current progress

1. Implemented a serial implementation of `pfind` using BFS
2. Implemented recursive and iterative `pfind` versions
3. Prepared a dataset suite for benchmarking
4. Performed basic performance analysis of iterative `pfind` version

Current progress

1. Implemented a serial implementation of pfind using BFS
2. Implemented recursive and iterative pfind versions
3. Prepared a dataset suite for benchmarking
4. Performed basic performance analysis of iterative pfind version

Work in progress

1. Implementing heuristics to allow more granular parallelism
 - Switching to priority queue to prioritize larger directories
2. Performing more thorough benchmarking with all competitors

Evaluation plan

Competitors for `pfind`

- Standard `find` that is used as baseline serial implementation
- MPI version of `find` from [1]

Scoring

1. Runtime
2. Resource usage

Evaluation plan: datasets

Datasets

- 1GB of Git repositories
 - “Awesome” list
 - Electron app framework
 - freeCodeCamp’s curriculum
 - JavaScript algorithm implementations
 - AMD MLIR AI engine toolchain
 - NVIDIA Linux GPU kernel module
 - Paper Minecraft server

Evaluation plan: datasets

- **Synthetic directory structures**
 - from high fan-out trees with few files to shallow trees with many files
- **Directories “in the wild”**
 - /home/ on PACE and our own Unix home directories

Evaluation plan: challenges

1. Varying Directory Sizes

- Large directories take too long to process
- Small directories don't show parallelization benefits

Evaluation plan: challenges

1. Varying Directory Sizes

- Large directories take too long to process
- Small directories don't show parallelization benefits

2. Ideal Directory Structure Hard to Find

- Structures showcasing parallel benefits are rare
- Need balance in depth, file count, and size

Evaluation plan: challenges

1. Varying Directory Sizes

- Large directories take too long to process
- Small directories don't show parallelization benefits

2. Ideal Directory Structure Hard to Find

- Structures showcasing parallel benefits are rare
- Need balance in depth, file count, and size

3. Modern vs. Older Hardware

- Default find is efficient on new hardware
- Older systems with slow I/O benefit more from parallelization

Benchmarking results

Git:

- `pfind` does better on real time, but uses more compute time.
- 1.8x to 4x speedup on real time, but 2x to 4x as much compute time

Benchmarking results

Git:

- `pfind` does better on real time, but uses more compute time.
- 1.8x to 4x speedup on real time, but 2x to 4x as much compute time

Synthetic:

- Similar performance between `pfind` and `find` (but more compute time)

Benchmarking results

Git:

- `pfind` does better on real time, but uses more compute time.
- 1.8x to 4x speedup on real time, but 2x to 4x as much compute time

Synthetic:

- Similar performance between `pfind` and `find` (but more compute time)

In the wild:

- `pfind` is usually about 2x faster, sometimes 2x slower

Conclusions

- The main bottleneck is I/O, so it's hard to get lots of speedup
- Faster real time/wall time comes at higher computation time cost
- Computing heuristics takes time and resources, which incurs overhead

Bibliography

- [1] E. Ong, E. Lusk, and W. Gropp, “Scalable Unix Commands for Parallel Processors: A High-Performance Implementation,” pp. 410–418, 2001.